

04 - Loading and Preparing Electric Consumption Data with Spark Core

Author: Adrián Romero Flores

Repository: [link](#)

Disclaimer: This document has been automatically generated from a Jupyter notebook and may contain formatting artifacts or code output inconsistencies. For the most accurate representation, refer to the original notebook in the repository.

Overview

1. Introduction
2. Environment Setup
3. Data Loading
4. Data Preparation
5. Results and Conclusions

Introduction

This project aims to load and prepare electric consumption data using Spark Core. The data is provided in CSV format and contains information about electric consumption in a specific region. The goal is to load the data into a Spark DataFrame, perform some transformations, and save the results in a format suitable for further analysis.

The document structure is as follows, in section Environment Setup we will talk about the execution environment. In section Data Loading, we will load the data into a Spark DataFrame and perform some initial transformations. In section Data Preparation, we will perform additional transformations to prepare the data for analysis. Finally, in section Results and Conclusions, we will summarize the results and provide conclusions based on the analysis.

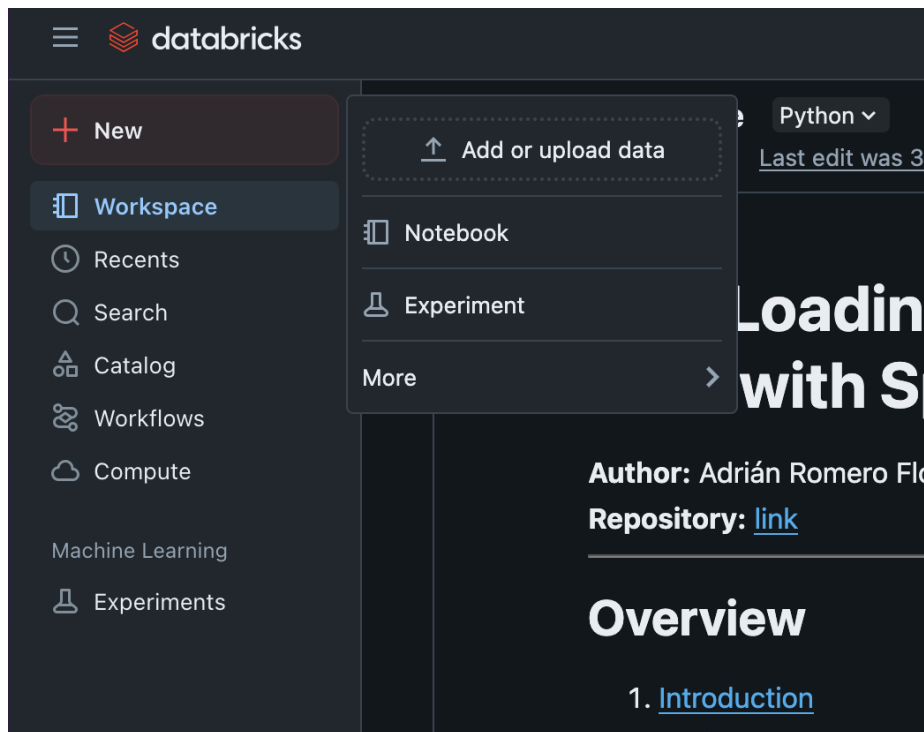
Environment Setup

For the execution of this notebook, we will use the DataBricks platform, which provides a cloud-based environment for running Spark applications. The notebook is written in Python and uses the PySpark library to interact with Spark. Here is a screenshot of the environment and the cluster details:

The image shows two screenshots of the Databricks interface. The top screenshot displays a notebook titled "04 - Loading and Preparing Electric Consumption Data with Spark Core" by Adrián Romero Flores. The notebook content includes an overview with links to Introduction, Environment Setup, Data Loading, Data Preparation, and Results and Conclusions. The introduction text states: "This project aims to load and prepare electric consumption data using Spark Core. The data is provided in CSV format and contains information about electric consumption in a specific region. The goal is to load the data into a Spark DataFrame, perform some transformations, and save the results in a format suitable for further analysis." The bottom screenshot shows the "My Cluster" configuration page, which includes settings for Databricks Runtime Version (12.2 LTS), Driver type (Community Optimized), Instance (15.3 GB Memory, 2 Cores), Spark (JDBC/ODBC), Spark config (spark.databricks.rocksdB.fileManager.useCommitService false), and Environment variables (PYSPARK_PYTHON=/databricks/python3/bin/python3).

Data Loading

First, we need to upload the CSV to the Databricks environment. We can do this by using the “Upload Data” button in the Data section of the workspace.



Once the file is uploaded, we can use the following code to load the data into a Spark DataFrame. The CSV file contains electric consumption data, and we will use the `spark.read.csv` method to read the file. We will also specify the schema of the DataFrame using the `schema` parameter. The schema is defined using the `StructType` and `StructField` from the `pyspark.sql.types` module.

```
from pyspark.sql.types import StructType, StructField, StringType, DoubleType
# define the struct of types
schema = StructType([
    StructField("IDENTIFICADOR", StringType(), True),
    StructField("ANOMES", StringType(), True),
    StructField("CNAE", StringType(), True),
    StructField("PRODUCTO", StringType(), True),
    StructField("MERCADO", StringType(), True),
    StructField("ACTIVA_H1", DoubleType(), True),
    StructField("ACTIVA_H2", DoubleType(), True),
    StructField("ACTIVA_H3", DoubleType(), True),
    StructField("ACTIVA_H4", DoubleType(), True),
    StructField("ACTIVA_H5", DoubleType(), True),
    StructField("ACTIVA_H6", DoubleType(), True),
    StructField("ACTIVA_H7", DoubleType(), True),
    StructField("ACTIVA_H8", DoubleType(), True),
```

```

    StructField("ACTIVA_H9", DoubleType(), True),
    StructField("ACTIVA_H10", DoubleType(), True),
    StructField("ACTIVA_H11", DoubleType(), True),
    StructField("ACTIVA_H12", DoubleType(), True),
    StructField("ACTIVA_H13", DoubleType(), True),
    StructField("ACTIVA_H14", DoubleType(), True),
    StructField("ACTIVA_H15", DoubleType(), True),
    StructField("ACTIVA_H16", DoubleType(), True),
    StructField("ACTIVA_H17", DoubleType(), True),
    StructField("ACTIVA_H18", DoubleType(), True),
    StructField("ACTIVA_H19", DoubleType(), True),
    StructField("ACTIVA_H20", DoubleType(), True),
    StructField("ACTIVA_H21", DoubleType(), True),
    StructField("ACTIVA_H22", DoubleType(), True),
    StructField("ACTIVA_H23", DoubleType(), True),
    StructField("ACTIVA_H24", DoubleType(), True),
    StructField("ACTIVA_H25", DoubleType(), True),
    StructField("REACTIVA_H1", DoubleType(), True),
    StructField("REACTIVA_H2", DoubleType(), True),
    StructField("REACTIVA_H3", DoubleType(), True),
    StructField("REACTIVA_H4", DoubleType(), True),
    StructField("REACTIVA_H5", DoubleType(), True),
    StructField("REACTIVA_H6", DoubleType(), True),
    StructField("REACTIVA_H7", DoubleType(), True),
    StructField("REACTIVA_H8", DoubleType(), True),
    StructField("REACTIVA_H9", DoubleType(), True),
    StructField("REACTIVA_H10", DoubleType(), True),
    StructField("REACTIVA_H11", DoubleType(), True),
    StructField("REACTIVA_H12", DoubleType(), True),
    StructField("REACTIVA_H13", DoubleType(), True),
    StructField("REACTIVA_H14", DoubleType(), True),
    StructField("REACTIVA_H15", DoubleType(), True),
    StructField("REACTIVA_H16", DoubleType(), True),
    StructField("REACTIVA_H17", DoubleType(), True),
    StructField("REACTIVA_H18", DoubleType(), True),
    StructField("REACTIVA_H19", DoubleType(), True),
    StructField("REACTIVA_H20", DoubleType(), True),
    StructField("REACTIVA_H21", DoubleType(), True),
    StructField("REACTIVA_H22", DoubleType(), True),
    StructField("REACTIVA_H23", DoubleType(), True),
    StructField("REACTIVA_H24", DoubleType(), True),
    StructField("REACTIVA_H25", DoubleType(), True)
  ])
# read csv without header
df = spark.read.csv("/FileStore/tables/endesaAgregada.csv", header=False, schema=schema)

```

To understand the structure and content of the dataframe, we can use the following code that displays the schema, the first 5 rows, the number of rows and columns, and summary statistics for each column.

```
# dataframe schema
for field in df.dtypes:
    print(f"column name: {field[0]}, type: {field[1]}")
# print 5 first rows
df.show(5)
# number of rows and columns
rows = df.count()
columns = len(df.columns)
print(f"number of rows: {rows}")
print(f"number of columns: {columns}")
# some stats -> the format is not so good for visualizing...
df.describe().show()

column name: IDENTIFICADOR, type: string
column name: ANOMES, type: string
column name: CNAE, type: string
column name: PRODUCTO, type: string
column name: MERCADO, type: string
column name: ACTIVA_H1, type: double
column name: ACTIVA_H2, type: double
column name: ACTIVA_H3, type: double
column name: ACTIVA_H4, type: double
column name: ACTIVA_H5, type: double
column name: ACTIVA_H6, type: double
column name: ACTIVA_H7, type: double
column name: ACTIVA_H8, type: double
column name: ACTIVA_H9, type: double
column name: ACTIVA_H10, type: double
column name: ACTIVA_H11, type: double
column name: ACTIVA_H12, type: double
column name: ACTIVA_H13, type: double
column name: ACTIVA_H14, type: double
column name: ACTIVA_H15, type: double
column name: ACTIVA_H16, type: double
column name: ACTIVA_H17, type: double
column name: ACTIVA_H18, type: double
column name: ACTIVA_H19, type: double
column name: ACTIVA_H20, type: double
column name: ACTIVA_H21, type: double
column name: ACTIVA_H22, type: double
column name: ACTIVA_H23, type: double
column name: ACTIVA_H24, type: double
column name: ACTIVA_H25, type: double
```

column name: REACTIVA_H1, type: double
column name: REACTIVA_H2, type: double
column name: REACTIVA_H3, type: double
column name: REACTIVA_H4, type: double
column name: REACTIVA_H5, type: double
column name: REACTIVA_H6, type: double
column name: REACTIVA_H7, type: double
column name: REACTIVA_H8, type: double
column name: REACTIVA_H9, type: double
column name: REACTIVA_H10, type: double
column name: REACTIVA_H11, type: double
column name: REACTIVA_H12, type: double
column name: REACTIVA_H13, type: double
column name: REACTIVA_H14, type: double
column name: REACTIVA_H15, type: double
column name: REACTIVA_H16, type: double
column name: REACTIVA_H17, type: double
column name: REACTIVA_H18, type: double
column name: REACTIVA_H19, type: double
column name: REACTIVA_H20, type: double
column name: REACTIVA_H21, type: double
column name: REACTIVA_H22, type: double
column name: REACTIVA_H23, type: double
column name: REACTIVA_H24, type: double
column name: REACTIVA_H25, type: double

IDENTIFICADOR	ANOMES	CNAE	PRODUCTO	MERCADO	ACTIVA_H1	ACTIVA_H2	ACTIVA_H3	ACTIVA_H4	ACTIVA_H5
32116	201507	T1	P1	M1	512.0	392.0	365.0	377.0	349.0
27062	201507	T1	P1	M1	0.0	0.0	0.0	0.0	0.0
86254	201507	T1	P38	M2	148.0	137.0	121.0	137.0	98.0
73313	201503	T1	P13	M2	318.0	272.0	149.0	148.0	136.0
46321	201411	T1	P1	M1	50.0	28.0	54.0	27.0	59.0

only showing top 5 rows

number of rows: 1173369

number of columns: 55

summary	IDENTIFICADOR	ANOMES	CNAE	PRODUCTO	MERCADO	ACTIVA_H1
count	1173369	1173369	1173369	1173369	1173369	1173369
mean	50087.853757002275	201488.48303133968	null	null	null	236.76462349013823
stddev	28912.129623636032	36.15787489052472	null	null	null	531.4217359903192
min	1	201309	T1	P1	M1	-6.5
max	99999	201510	T2	P98	M2	75096.0

```

+-----+-----+-----+-----+-----+-----+-----+-----+

```

As we can see in the output, the DataFrame contains negative values in the ACTIVA columns, which indicates that there are some errors in the data. We will need to handle these errors in the next section.

Data Preparation

First of all, let's remove the columns that are not needed for the analysis. We can do this using the `drop` method of the DataFrame.

```

# let's remove the ACTIVA_H25 column
df = df.drop("ACTIVA_H25")
# let's remove all REACTIVA_HX columns
reactiva_columns = [f"REACTIVA_H{i}" for i in range(1, 25)]

df = df.drop(*reactiva_columns)
print(f"columns after removal: {len(df.columns)} columns")

columns after removal: 30 columns

```

Once we have removed the unnecessary columns, we can proceed to handle the null values. We can use the `isnan` and `col(x).isNull()` functions to check for null values in the DataFrame. In this case there wasn't any null or nan values in the dataset, so we can proceed to the next step.

```

from pyspark.sql.functions import col, isnan, when, count

# there is any null value?
null_summary = df.select([
    count(when(col(c).isNull() | isnan(c), c)).alias(c) for c in df.columns
])

null_summary.show()

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+
|IDENTIFICADOR|ANOMES|CNAE|PRODUCTO|MERCADO|ACTIVA_H1|ACTIVA_H2|ACTIVA_H3|ACTIVA_H4|ACTIVA_H5|
+-----+-----+-----+-----+-----+-----+-----+-----+
|              |0|    |0|    |0|    |0|    |0|    |0|    |0|    |0|
+-----+-----+-----+-----+-----+-----+-----+-----+

```

Finally, we are going to remove the negative values in the ACTIVA columns. We can do this using the `filter` method of the DataFrame. We will filter out the rows where the ACTIVA columns are less than 0. This will ensure that we only keep the rows with valid values.

```

# is there any negative value in activa columns?
df_negatives = df.filter(
    " OR ".join([f"{columna} < 0" for columna in activa_columns])
)

```

```
print(f"number of rows with negative values: {df_filtered.count()}")
df = df.subtract(df_negatives)
```

```
df.describe().show()
```

```
number of rows with negative values: 4
```

summary	IDENTIFICADOR	ANOMES	CNAE	PRODUCTO	MERCADO	ACTIVA_H1
count	1173365	1173365	1173365	1173365	1173365	1173365
mean	50087.835313819654	201488.4831386653	null	null	null	236.76357399445186
stddev	28912.143205991142	36.157790023692854	null	null	null	531.4213745115978
min	1	201309	T1	P1	M1	0.0
max	99999	201510	T2	P98	M2	75096.0

Now, let's visualize some data, for example, we can view the aggregated consumption per hour grouped by cnae. We can use the `groupBy` method of the `DataFrame` to group the data by `CNAE`, and then use the `agg` method to calculate the sum of the `ACTIVA` columns. We will also use the `matplotlib` library to plot the data.

```
from pyspark.sql.functions import sum as _sum
import matplotlib.pyplot as plt

activa_columns = [f"ACTIVA_H{i}" for i in range(1, 25)]
agg_cnae = df.groupBy("CNAE").agg(*[_sum(col).alias(col) for col in activa_columns])

agg_cnae_pd = agg_cnae.toPandas()
agg_cnae_pd.set_index("CNAE", inplace=True)
agg_cnae_pd = agg_cnae_pd.transpose()
agg_cnae_pd.index = [f"H{i}" for i in range(1, 25)]

plt.figure(figsize=(10, 6))
for cnae_type in agg_cnae_pd.columns:
    plt.plot(agg_cnae_pd.index, agg_cnae_pd[cnae_type], marker='o', label=f'CNAE {cnae_type}')

plt.title("Evolución del Consumo Activo por Hora según Tipo CNAE")
plt.xlabel("Hora del Día")
plt.ylabel("Consumo Activo Total")
plt.legend()
plt.grid(True)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```


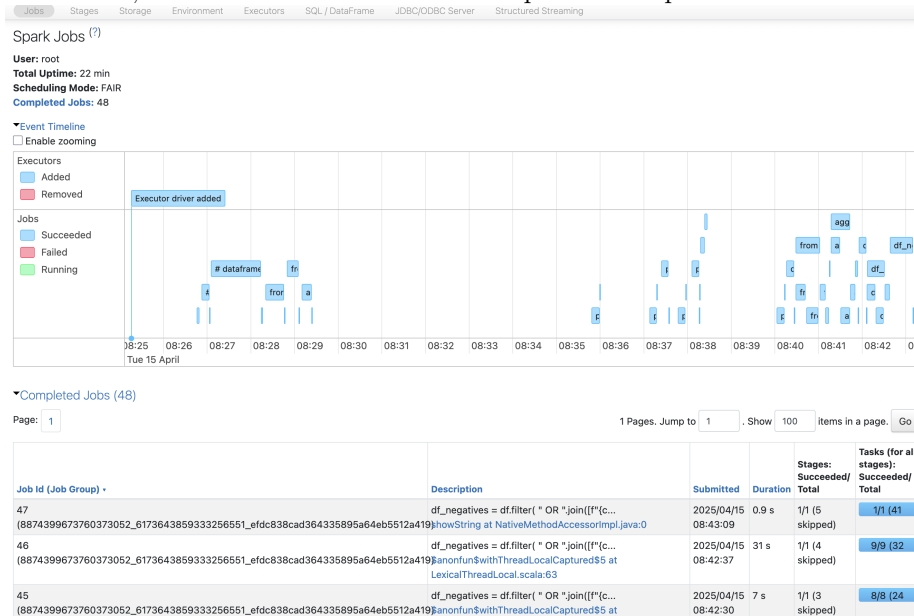



Figure 1: Graph ACTIVA grouped by CNAE

Monitoring the Spark Jobs

We can monitor the Spark jobs using the Spark UI that Databricks provide. The Spark UI provides a web interface for monitoring and debugging Spark applications. It shows the status of the jobs, stages, and tasks, as well as the execution plan and performance metrics.



Results and Conclusions

Based on the results of the analysis, we can conclude that the electric consumption data contains some errors, such as negative values in the **ACTIVA** columns. We have successfully loaded and prepared the data for further analysis by removing unnecessary columns, handling null values, and filtering out invalid values. The data is now ready for further analysis and visualization.

The visualization of aggregated hourly consumption segmented by CNAE, reveals discernible trends and consumption patterns. In the case of CNAE, the profile T1 demonstrates a consumption pattern typical of residential or human usage: minimal activity during nocturnal hours, with notable increases during the late afternoon and evening—corresponding with common periods of occupancy and usage. Conversely, the T2 profile reflects a pattern indicative of industrial or commercial activity, characterized by a peak in the morning hours followed by a significant decline during midday, aligning with standard lunch breaks.